

Using Neural Network Potentials to Predict Molecular Energies of Small Organic Molecules

Marco Huang, BioE 142 Final project

1 . Introduction

To calculate the potential energy of a molecule from first principles, the Schrödinger Equation must be solved. The density functional theory (DFT), a quantum mechanical approach, can provide accurate energy calculations [1]. As the number of atoms grows, though, the process becomes computationally costly. Neural network potentials (NNPs) allow for accuracies close to what a quantum mechanical approach would produce while requiring a fraction of the computational cost [2].

In this paper, an Accurate Neural network Engine for Molecular Energies (ANI) framework was implemented and evaluated to predict the molecular energies for molecules with up to four heavy atoms [3]. ANIs were originally developed by Smith et al. in 2017 [2,3]. These frameworks allow for the training of element specific atomic neural networks using atomic environment vectors (AEV) [2,3]. The environment around each atom is encoded into a numerical descriptor with a fixed length. These vectors capture radial and angular features to ensure rotational, translational, and permutational invariance. The way the NNP works for predicting molecular energies is as follows. For each element type (H,C,N,O), there is a separate neural network to predict energy contributions. The molecular energy is the sum of these per atom energy contributions. The act of decomposing molecules into their atoms allows the model to generalize its understanding on different sized molecules with various compositions. Smith et al. trained their ANI on approximately 17 million conformations from 57,951 molecules that contained up to 8 heavy atoms (H, C, N, and O). Their model was able to accurately predict molecular energies of molecules much larger than those in the training dataset.

The ANI framework was used in this paper to train neural network potentials on a subset of the GDB database. The subset contained molecules up to 4 heavy atoms. The goal of this study was to investigate how to tune the model architecture and hyperparameters to minimize error in prediction. The standard threshold for proper predictions on chemical behavior is an mean absolute error (MAE) at or under 1kcal/mol. The best model came from the fifth fold of a 10-fold cross validation. Its MAE = 1.016 kcal/mol and its RMSE was 1.38 kcal/mol. The cross fold validation mean was 1.675 kcal/mol. With more data, more availability in computational resources, and longer training runs with more epochs, it can be anticipated that the MAE can go even lower.

2 . Methods

The ANI_GDB_s01_to_s04 dataset was used for this study. It contains DFT computed energies for molecular conformations containing up to 4 heavy atoms. These molecules are composed of H, C, N, and O. The dataset contains 864,896 total conformations. 80% (691,918) of the data went into training, 10% (86,489) went into validation, and the remaining 10% (86,489) went into a held out test set.

AEVs were computed via the TorchANI library [3]. The AEV was configured to have a radial cutoff (rcr) of 5.2 angstroms and the angular cutoff (rca) of 3.5 angstroms. There are 64 radial elements (4 atom types * 16 shifts) and 320 angular elements (10 unique atom pairs * 4 radial shifts * 8 angular shifts). Thus the length of the AEV fed into each network is 384.

An atomic neural network was instantiated for each heavy atom and trained jointly with the TorchANI ANIModel Class. The neural network architecture for each heavy atom type is as follows: 384 input nodes to 512 hidden nodes, to 256 hidden nodes, to 128 hidden nodes to 1 output node. Each individual atomic network had 363,265 parameters. Combining the four networks together, there are 1,453,060 parameters. All hidden layers used Batch Norm 1d with a ReLU activation function that follows it. Various dropout rates at 0, 0.05, and 0.15 were experimented with. This was due to the result that the model did not seem to be overfitting during training and validation. Thus, dropout was lowered to explore if loss could continue dropping.

The Adam Optimizer with mean squared error loss was implemented to train this model. A mini-batch size of 8,192 was used. 8,192 was the default batch size set by the notebook template and was not changed during training. This was set due to the result that the training loss decreased at a smooth rate and the fact that training on the full dataset took roughly 10 minutes. A ReduceLROnPlateau learning rate scheduler was implemented as well. There are hyperparameters to set for this scheduler: the learning rate factor and the patience. After validation loss did not improve for 10 epochs, the learning rate would be multiplied by factor = 0.5. ReduceLROnPlateau also allows the user to set a learning rate floor to avoid making the learning rate so small that the steps towards the minimum are negligible. The learning rate floor for this study was set to 1e-8.

Adjustments to the learning rate, number of epochs, and L2 regression hyperparameters were made to narrow down which values would produce the lowest MAE in kcal/mol. The best model after hyperparameter tuning came out to produce an MAE of 1.17 kcal/mol. It was trained on a learning rate of 1e-4 and 200 epochs. Throughout all experiments with hyperparameter tuning, an L2 regularization hyperparameter was set to 1e-5. This was the initial L2 hyperparameter that was set for initial exploration. L2 regression penalizes larger weights and increasing the

hyperparameter would help if the model was overfitting. Throughout training, however, signs of overfitting did not appear, so increasing this hyperparameter would hurt the result.

Taking the best model configurations (hyperparameter tuning, model architecture engineering, and regularization strategies), a 10-fold cross validation was configured to explore model performance on different portions of the data and to understand which sections of data produced the lowest error. 8 of the 10 portions were used for training, 1 for validation, and 1 for a separate held-out testing set. Instead of having 1 test set, the remaining 2 portions were split into 2. The validation set was used for early stopping and the learning rate scheduler to keep track of loss.

3 . Results

Hidden Layers	Learning Rate	Num Epochs	L2 Regularization	Dropout	Patience	Test MAE (kcal/mol)
[384, 256, 192, 128, 1]	1e-04	30	1e-05	0.15	10	7.89
[384, 256, 192, 128, 1]	1e-08	100	1e-05	0.15	10	349.59
[384, 256, 192, 128, 1]	1e-05	100	1e-05	0.15	10	7.57
[384, 256, 192, 128, 1]	1e-04	100	1e-05	0.15	10	3.51
[384, 256, 192, 128, 1]	1e-04	200	1e-05	0.15	10	2.42
[384, 512, 256, 128, 1]	1e-04	200	1e-05	0.15	10	1.81
[384, 512, 256, 128, 1]	1e-04	200	1e-05	0.05	10	1.17
[384, 512, 256, 128, 1]	1e-04	200	1e-05	0	10	4.92
[384, 512, 256, 128, 1]	1e-04	200	1e-05	0.05	5	2.68

Table 1: table outlining the hyperparameter configurations tested and their resulting test MAE in kcal/mol.

The initial configuration with the structure [384, 256, 192, 128, 1], a learning rate of 1e-4, 30 epochs, and a drop out rate of 0.15 achieved an MAE of 7.89 kcal/mol. At 30 epochs, the loss curves have not converged yet, confirming that the model needed more training and more epochs.

Raising the number of epochs to 100, several initial learning rates were tested and critical effects were noticed. Lowering the initial learning rate to 1e-8 caused the model to produce an MAE of 349.59 kcal/mol. With the same hyperparameters and tuning learning rate to 1e-5, the MAE was reduced to 7.57 kcal/mol. At these learning rates, the progress towards finding the minimum loss is negligible because parameter updates became too small to make any meaningful progress. Increasing the learning rate of 1e-4 produced an MAE of 3.51 kcal/mol and produced loss curves

with the fewest fluctuations. Keeping the same parameters, but using double the number of epochs produced an MAE of 2.42 kcal/mol

Because the loss curves continued to show signs that the model needed more training to converge, the neural network architecture was also engineered to have a higher capacity [384, 512, 256, 128, 1]. Across the four models for each atom type, there are 1,453,060. This architecture contains nearly twice as many parameters as the initial architecture which had 695,556 parameters. Nearly doubling the number of parameters reduced the MAE to 1.81 kcal/mol.

At this point, the training loss curve slope began to decrease. There were still no signs of overfitting. Therefore, the number of epochs was raised to 200 and the drop out rate was lowered to 0.05. This model achieved an MAE of 1.17 kcal/mol, the lowest of all trials. The training loss curve's slope was nearly 0 at 200 epochs. To see if drop out was preventing the loss curve from dropping even further, it was completely removed. The MAE spiked to 4.92 kcal/mol. The learning rate's patience was also reduced to 5 to verify that it was not the factor that makes the model converge completely. This edit made the model produce an MAE of 2.68 kcal/mol, but prematurely flattened the training loss curve at 100 epochs.




Figure 1: Training and Loss Curves for the best model produced while model tuning(left). Predicted vs Ground truth of said model (right)

Fold	Fold test MAE (kcal/mol)	Fold test RMSE (kcal/mol)	Original test set MAE (kcal/mol)	Original test set RMSE (kcal/mol)
1	1.477	1.893	1.47	1.88
2	1.329	1.656	1.32	1.64
3	2.132	2.647	2.13	2.66
4	1.929	2.433	1.93	2.43
5	1.016	1.393	1.01	1.38
6	2.620	3.087	2.61	3.07
7	1.181	1.676	1.17	1.67
8	2.011	2.522	1.99	2.48
9	1.078	1.519	1.07	1.52
10	1.973	2.388	1.96	2.36
Cross Fold Validation Error Mean and Standard deviation	1.675 +/- 0.506 kcal/mol	2.123 +/- 0.539 kcal/mol	1.666 +/- 0.533 kcal/mol	2.109 +/- 0.565 kcal/mol

Table 2: 10-fold cross validation errors on fold test set and original test set.

Fold five produced the lowest errors for each category of error and was further broken down to error by element composition.




Figure 2: Training and Loss Curves for the best fold (fold 5) of the 10 fold cross validation (left). Predicted vs Ground truth of fold 5.

Element Composition	Count Data Points	MAE (kcal/mol)	RMSE (kcal/mol)
H + C	17,639	0.84	1.24
H + C + N	32,299	1.13	1.52
H + C + O	25,244	0.84	1.14
H + C + N + O	11,307	1.28	1.64

Table 3: Error by element composition

Without nitrogen, the MAE for both H + C and H + C + O was 0.84 kcal/mol. With nitrogen, the MAE increased to beyond 1.




Figure 3: Predicted vs Ground truth of fold 5 by element composition

Comparing H + C + N and H + C + N + O, one may notice that H + C + N + O seems more tightly clustered around the predicted equals ground truth than H + C + N even though H + C +

N has a lower RMSE. This could be because there are many data points, and the few that are not so tightly clustered together give the illusion that the density around the line is lower.

4 . Conclusions

The ANI neural network can achieve near chemical accuracy energy predictions. This is demonstrated through its ability to predict near ground truth molecular energies even when training on a small subset of the GDB database containing molecules with up to 4 heavy atoms. Through tuning of the model, it was found that increasing the network width, lowering drop out rates, and keeping learning rates from being too small significantly impacted the model's performance. Specifically, a wider architecture with light drop out consistently outperformed models with narrower architectures.

Looking at the Predicted vs Ground truth of fold 5 (figure 3), we see that the predictions are tightly clustered around the predicted equals ground truth line. This suggests that the model does not systematically over or under predict the molecular energy.

When comparing the RMSE of fold 5 (1.38 kcal/mol) to the RMSE in the paper (1.3 kcal/mol), one can notice that the RMSE in this paper is 0.08 kcal/mol higher than that of the one in the paper. This can be due to several factors. The dataset used for this study contains 864,896 total conformations and up to 4 heavy atoms. The dataset that Smith et al. used contained roughly 17.2 million conformations and up to 8 heavy atoms. The dataset used for this study was roughly 5% of the total size of the dataset used by Smith et al. The AEV resolution for this study was also lower. With fewer molecules, the model is exposed to fewer types of bonds, functional groups, and structures during training. Thus having less information to predict energies of new molecules. Smith et al. also used a gaussian activation function with an exponential cost function. They mentioned that this approach gave 2-4x lower error than ReLU activation function and MSe loss function approach used in this study.

Future studies can look into potentially training with a larger subset of the dataset with more heavy atoms. Since the best models produced in this study did not show signs of overfitting, a potential path of research would be to explore even wider hidden layers in the neural network. A larger, more descriptive AEV would also provide the model with more training knowledge and could perhaps drop the error even further. Smith et al. used an AEV of length 768 which gave double the resolution than the 384 length AEV used in this model. One could experiment to see if a wider network with longer AEVs could potentially decrease MAE and RMSE even further. In terms of practicality, incorporating force predictions could potentially enable understanding of molecular dynamics. Overall, the artificial neural networks trained in this study were able to predict the molecular energies of thousands of molecular conformations to near chemical accuracy. One can expand on the findings to potentially achieve better results and practicality.

5 . References

1. W. Kohn and L. J. Sham, "Self-consistent equations including exchange and correlation effects," *Phys. Rev.*, 1965, **140**, A1133–A1138.
2. J. Behler and M. Parrinello, "Generalized neural-network representation of high-dimensional potential-energy surfaces," *Phys. Rev. Lett.*, 2007, **98**, 146401.
3. J. S. Smith, O. Isayev, and A. E. Roitberg, "ANI-1: an extensible neural network potential with DFT accuracy at force field computational cost," *Chem. Sci.*, 2017, **8**, 3192–3203.
4. X. Gao, F. Ramezanghorbani, O. Isayev, J. S. Smith, and A. E. Roitberg, "TorchANI: A free and open source PyTorch-based deep learning implementation of the ANI neural network potentials," *J. Chem. Inf. Model.*, 2020, **60**, 3408–3415.