

# Using Neural Network Potentials to Predict Molecular Energies of Small Organic Molecules

*Marco Huang, BioE 142 Final project*

## 1 . Introduction

To calculate the potential energy of a molecule from first principles, the Schrödinger Equation must be solved. The density functional theory (DFT), a quantum mechanical approach, can provide accurate energy calculations [1]. As the number of atoms grows, though, the process becomes computationally costly. Neural network potentials (NNPs) allow for accuracies close to what a quantum mechanical approach would produce while requiring a fraction of the computational cost [2].

In this paper, an Accurate Neural network Engine for Molecular Energies (ANI) framework was implemented and evaluated to predict the molecular energies for molecules with up to four heavy atoms [3]. ANIs were originally developed by Smith et al. in 2017 [2,3]. These frameworks allow for the training of element specific atomic neural networks using atomic environment vectors (AEV) [2,3]. The environment around each atom is encoded into a numerical descriptor with a fixed length. These vectors capture radial and angular features to ensure rotational, translational, and permutational invariance. The way the NNP works for predicting molecular energies is as follows. For each element type (H,C,N,O), there is a separate neural network to predict energy contributions. The molecular energy is the sum of these per atom energy contributions. The act of decomposing molecules into their atoms allows the model to generalize its understanding on different sized molecules with various compositions. Smith et al. trained their ANI on approximately 17 million conformations from 57,951 molecules that contained up to 8 heavy atoms (H, C, N, and O). Their model was able to accurately predict molecular energies of molecules much larger than those in the training dataset.

The ANI framework was used in this paper to train neural network potentials on a subset of the GDB database. The subset contained molecules up to 4 heavy atoms. The goal of this study was to investigate how to tune the model architecture and hyperparameters to minimize error in prediction. The standard threshold for proper predictions on chemical behavior is an mean absolute error (MAE) at or under 1kcal/mol. The best model came from the fifth fold of a 10-fold cross validation. Its MAE = 1.016 kcal/mol and its RMSE was 1.38 kcal/mol. The cross fold validation mean was 1.675 kcal/mol. With more data, more availability in computational resources, and longer training runs with more epochs, it can be anticipated that the MAE can go even lower.

## 2 . Methods

The ANI\_GDB\_s01\_to\_s04 dataset was used for this study. It contains DFT computed energies for molecular conformations containing up to 4 heavy atoms. These molecules are composed of H, C, N, and O. The dataset contains 864,896 total conformations. 80% (691,918) of the data went into training, 10% (86,489) went into validation, and the remaining 10% (86,489) went into a held out test set.

AEVs were computed via the TorchANI library [3]. The AEV was configured to have a radial cutoff (rcr) of 5.2 angstroms and the angular cutoff (rca) of 3.5 angstroms. There are 64 radial elements (4 atom types \* 16 shifts) and 320 angular elements (10 unique atom pairs \* 4 radial shifts \* 8 angular shifts). Thus the length of the AEV fed into each network is 384.

An atomic neural network was instantiated for each heavy atom and trained jointly with the TorchANI ANIModel Class. The neural network architecture for each heavy atom type is as follows: 384 input nodes to 512 hidden nodes, to 256 hidden nodes, to 128 hidden nodes to 1 output node. Each individual atomic network had 363,265 parameters. Combining the four networks together, there are 1,453,060 parameters. All hidden layers used Batch Norm 1d with a ReLU activation function that follows it. Various dropout rates at 0, 0.05, and 0.15 were experimented with. This was due to the result that the model did not seem to be overfitting during training and validation. Thus, dropout was lowered to explore if loss could continue dropping.

The Adam Optimizer with mean squared error loss was implemented to train this model. A mini-batch size of 8,192 was used. 8,192 was the default batch size set by the notebook template and was not changed during training. This was set due to the result that the training loss decreased at a smooth rate and the fact that training on the full dataset took roughly 10 minutes. A ReduceLROnPlateau learning rate scheduler was implemented as well. There are hyperparameters to set for this scheduler: the learning rate factor and the patience. After validation loss did not improve for 10 epochs, the learning rate would be multiplied by factor = 0.5. ReduceLROnPlateau also allows the user to set a learning rate floor to avoid making the learning rate so small that the steps towards the minimum are negligible. The learning rate floor for this study was set to 1e-8.

Adjustments to the learning rate, number of epochs, and L2 regression hyperparameters were made to narrow down which values would produce the lowest MAE in kcal/mol. The best model after hyperparameter tuning came out to produce an MAE of 1.17 kcal/mol. It was trained on a learning rate of 1e-4 and 200 epochs. Throughout all experiments with hyperparameter tuning, an L2 regularization hyperparameter was set to 1e-5. This was the initial L2 hyperparameter that was set for initial exploration. L2 regression penalizes larger weights and increasing the

hyperparameter would help if the model was overfitting. Throughout training, however, signs of overfitting did not appear, so increasing this hyperparameter would hurt the result.

Taking the model configurations that produced MAE 1.17 kcal/mol (hyperparameter tuning, model architecture engineering, and regularization strategies), a 10-fold cross validation was configured to explore model performance on different portions of the data and to understand which sections of data produced the lowest error. 8 of the 10 portions were used for training, 1 for validation, and 1 for a separate held-out testing set. Instead of having 1 test set, the remaining 2 portions were split into 2. The validation set was used for early stopping and the learning rate scheduler to keep track of loss.

### 3 . Results

| Hidden Layers                  | Learning Rate | Num Epochs | L2 Regularization | Dropout     | Patience  | Test MAE (kcal/mol) |
|--------------------------------|---------------|------------|-------------------|-------------|-----------|---------------------|
| [384, 256, 192, 128, 1]        | 1e-04         | 30         | 1e-05             | 0.15        | 10        | 7.89                |
| [384, 256, 192, 128, 1]        | 1e-08         | 100        | 1e-05             | 0.15        | 10        | 349.59              |
| [384, 256, 192, 128, 1]        | 1e-05         | 100        | 1e-05             | 0.15        | 10        | 7.57                |
| [384, 256, 192, 128, 1]        | 1e-04         | 100        | 1e-05             | 0.15        | 10        | 3.51                |
| [384, 256, 192, 128, 1]        | 1e-04         | 200        | 1e-05             | 0.15        | 10        | 2.42                |
| [384, 512, 256, 128, 1]        | 1e-04         | 200        | 1e-05             | 0.15        | 10        | 1.81                |
| <b>[384, 512, 256, 128, 1]</b> | <b>1e-04</b>  | <b>200</b> | <b>1e-05</b>      | <b>0.05</b> | <b>10</b> | <b>1.17</b>         |
| [384, 512, 256, 128, 1]        | 1e-04         | 200        | 1e-05             | 0           | 10        | 4.92                |
| [384, 512, 256, 128, 1]        | 1e-04         | 200        | 1e-05             | 0.05        | 5         | 2.68                |

**Table 1:** table outlining the hyperparameter configurations tested and their resulting test MAE in kcal/mol.


The initial configuration with the structure [384, 256, 192, 128, 1], a learning rate of 1e-4, 30 epochs, and a drop out rate of 0.15 achieved an MAE of 7.89 kcal/mol. At 30 epochs, the loss curves have not converged yet, confirming that the model needed more training and more epochs.

Raising the number of epochs to 100, several initial learning rates were tested and critical effects were noticed. Lowering the initial learning rate to 1e-8 caused the model to produce an MAE of 349.59 kcal/mol. With the same hyperparameters and tuning learning rate to 1e-5, the MAE was reduced to 7.57 kcal/mol. At these learning rates, the progress towards finding the minimum loss is negligible because parameter updates became too small to make any meaningful progress. Increasing the learning rate of 1e-4 produced an MAE of 3.51 kcal/mol and produced loss curves

with the fewest fluctuations. Keeping the same parameters, but using double the number of epochs produced an MAE of 2.42 kcal/mol

Because the loss curves continued to show signs that the model needed more training to converge, the neural network architecture was also engineered to have a higher capacity [384, 512, 256, 128, 1]. Across the four models for each atom type, there are 1,453,060. This architecture contains nearly twice as many parameters as the initial architecture which had 695,556 parameters. Nearly doubling the number of parameters reduced the MAE to 1.81 kcal/mol.

At this point, the training loss curve slope began to decrease. There were still no signs of overfitting. Therefore, the number of epochs was raised to 200 and the drop out rate was lowered to 0.05. This model achieved an MAE of 1.17 kcal/mol, the lowest of all trials. The training loss curve's slope was nearly 0 at 200 epochs. To see if drop out was preventing the loss curve from dropping even further, it was completely removed. The MAE spiked to 4.92 kcal/mol. The learning rate's patience was also reduced to 5 to verify that it was not the factor that makes the model converge completely. This edit made the model produce an MAE of 2.68 kcal/mol, but prematurely flattened the training loss curve at 100 epochs.




**Figure 1:** Training and Loss Curves for the best model produced while model tuning(left). Predicted vs Ground truth of said model (right)

| Fold                                                    | Fold test MAE (kcal/mol) | Fold test RMSE (kcal/mol) | Original test set MAE (kcal/mol) | Original test set RMSE (kcal/mol) |
|---------------------------------------------------------|--------------------------|---------------------------|----------------------------------|-----------------------------------|
| 1                                                       | 1.477                    | 1.893                     | 1.47                             | 1.88                              |
| 2                                                       | 1.329                    | 1.656                     | 1.32                             | 1.64                              |
| 3                                                       | 2.132                    | 2.647                     | 2.13                             | 2.66                              |
| 4                                                       | 1.929                    | 2.433                     | 1.93                             | 2.43                              |
| <b>5</b>                                                | <b>1.016</b>             | <b>1.393</b>              | <b>1.01</b>                      | <b>1.38</b>                       |
| 6                                                       | 2.620                    | 3.087                     | 2.61                             | 3.07                              |
| 7                                                       | 1.181                    | 1.676                     | 1.17                             | 1.67                              |
| 8                                                       | 2.011                    | 2.522                     | 1.99                             | 2.48                              |
| 9                                                       | 1.078                    | 1.519                     | 1.07                             | 1.52                              |
| 10                                                      | 1.973                    | 2.388                     | 1.96                             | 2.36                              |
| Cross Fold Validation Error Mean and Standard deviation | 1.675 +/- 0.506 kcal/mol | 2.123 +/- 0.539 kcal/mol  | 1.666 +/- 0.533 kcal/mol         | 2.109 +/- 0.565 kcal/mol          |

**Table 2:** 10-fold cross validation errors on fold test set and original test set.

Fold five produced the lowest errors for each category of error and was further broken down to error by element composition.




**Figure 2:** Training and Loss Curves for the best fold (fold 5) of the 10 fold cross validation (left). Predicted vs Ground truth of fold 5.

| Element Composition | Count Data Points | MAE (kcal/mol) | RMSE (kcal/mol) |
|---------------------|-------------------|----------------|-----------------|
| H + C               | 17,639            | 0.84           | 1.24            |
| H + C + N           | 32,299            | 1.13           | 1.52            |
| H + C + O           | 25,244            | 0.84           | 1.14            |
| H + C + N + O       | 11,307            | 1.28           | 1.64            |

**Table 3:** Error by element composition

Without nitrogen, the MAE for both H + C and H + C + O was 0.84 kcal/mol. With nitrogen, the MAE increased to beyond 1.



**Figure 3:** Predicted vs Ground truth of fold 5 by element composition

Comparing H + C + N and H + C + N + O, one may notice that H + C + N + O seems more tightly clustered around the predicted equals ground truth than H + C + N even though H + C +

N has a lower RMSE. This could be because there are many data points, and the few that are not so tightly clustered together give the illusion that the density around the line is lower.

## 4 . Conclusions

The ANI neural network can achieve near chemical accuracy energy predictions. This is demonstrated through its ability to predict near ground truth molecular energies even when training on a small subset of the GDB database containing molecules with up to 4 heavy atoms. Through tuning of the model, it was found that increasing the network width, lowering drop out rates, and keeping learning rates from being too small significantly impacted the model's performance. Specifically, a wider architecture with light drop out consistently outperformed models with narrower architectures.

Looking at the Predicted vs Ground truth of fold 5 (figure 3), we see that the predictions are tightly clustered around the predicted equals ground truth line. This suggests that the model does not systematically over or under predict the molecular energy.

When comparing the RMSE of fold 5 (1.38 kcal/mol) to the RMSE in the paper (1.3 kcal/mol), one can notice that the RMSE in this paper is 0.08 kcal/mol higher than that of the one in the paper. This can be due to several factors. The dataset used for this study contains 864,896 total conformations and up to 4 heavy atoms. The dataset that Smith et al. used contained roughly 17.2 million conformations and up to 8 heavy atoms. The dataset used for this study was roughly 5% of the total size of the dataset used by Smith et al. The AEV resolution for this study was also lower. With fewer molecules, the model is exposed to fewer types of bonds, functional groups, and structures during training. Thus having less information to predict energies of new molecules. Smith et al. also used a gaussian activation function with an exponential cost function. They mentioned that this approach gave 2-4x lower error than ReLU activation function and MSe loss function approach used in this study.

Future studies can look into potentially training with a larger subset of the dataset with more heavy atoms. Since the best models produced in this study did not show signs of overfitting, a potential path of research would be to explore even wider hidden layers in the neural network. A larger, more descriptive AEV would also provide the model with more training knowledge and could perhaps drop the error even further. Smith et al. used an AEV of length 768 which gave double the resolution than the 384 length AEV used in this model. One could experiment to see if a wider network with longer AEVs could potentially decrease MAE and RMSE even further. In terms of practicality, incorporating force predictions could potentially enable understanding of molecular dynamics. Overall, the artificial neural networks trained in this study were able to predict the molecular energies of thousands of molecular conformations to near chemical accuracy. One can expand on the findings to potentially achieve better results and practicality.

## 5 . References

1. W. Kohn and L. J. Sham, "Self-consistent equations including exchange and correlation effects," *Phys. Rev.*, 1965, **140**, A1133–A1138.
2. J. Behler and M. Parrinello, "Generalized neural-network representation of high-dimensional potential-energy surfaces," *Phys. Rev. Lett.*, 2007, **98**, 146401.
3. J. S. Smith, O. Isayev, and A. E. Roitberg, "ANI-1: an extensible neural network potential with DFT accuracy at force field computational cost," *Chem. Sci.*, 2017, **8**, 3192–3203.
4. X. Gao, F. Ramezanghorbani, O. Isayev, J. S. Smith, and A. E. Roitberg, "TorchANI: A free and open source PyTorch-based deep learning implementation of the ANI neural network potentials," *J. Chem. Inf. Model.*, 2020, **60**, 3408–3415.



# Final Project

This notebook is adapted from here: [https://aiqm.github.io/torchani/examples/nnp\\_training.html](https://aiqm.github.io/torchani/examples/nnp_training.html)

## Data preparation

1. Create a working directory: `/global/scratch/users/[USER_NAME]/[DIR_NAME]` . Replace the [USER\_NAME] with yours and specify a [DIR\_NAME] you like.
2. Copy the Jupyter Notebook to the working directory
3. Download the ANI dataset `ani_dataset_gdb_s01_to_s04.h5` from bCourses and upload it to the working directory

```
In [1]: import numpy as np
        from tqdm import tqdm
        import torch
        import torch.nn as nn
        import torchani
        import matplotlib.pyplot as plt
```

## Use GPU

```
In [2]: device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
        print(device)
```

cuda

## Set up AEV computer

AEV: Atomic Environment Vector (atomic features)

Ref: Chem. Sci., 2017, 8, 3192

```
In [3]: def init_aev_computer():
        Rcr = 5.2
        Rca = 3.5
        EtaR = torch.tensor([16], dtype=torch.float, device=device)
        ShfR = torch.tensor([
            0.900000, 1.168750, 1.437500, 1.706250,
            1.975000, 2.243750, 2.512500, 2.781250,
            3.050000, 3.318750, 3.587500, 3.856250,
```

```

    4.125000, 4.393750, 4.662500, 4.931250
], dtype=torch.float, device=device)

EtaA = torch.tensor([8], dtype=torch.float, device=device)
Zeta = torch.tensor([32], dtype=torch.float, device=device)
ShfA = torch.tensor([0.90, 1.55, 2.20, 2.85], dtype=torch.float, device=device)
ShfZ = torch.tensor([
    0.19634954, 0.58904862, 0.9817477, 1.37444680,
    1.76714590, 2.15984490, 2.5525440, 2.94524300
], dtype=torch.float, device=device)

num_species = 4
aev_computer = torchani.AEVComputer(
    Rcr, Rca, EtaR, ShfR, EtaA, Zeta, ShfA, ShfZ, num_species
)
return aev_computer

aev_computer = init_aev_computer()
aev_dim = aev_computer.aev_length
print(aev_dim)

```

384

## Prepare dataset & split

```

In [4]: def load_ani_dataset(dspath):
    self_energies = torch.tensor([
        0.500607632585, -37.8302333826,
        -54.5680045287, -75.0362229210
    ], dtype=torch.float, device=device)
    energy_shifter = torchani.utils.EnergyShifter(None)
    species_order = ['H', 'C', 'N', 'O']

    dataset = torchani.data.load(dspath)
    dataset = dataset.subtract_self_energies(energy_shifter, species_order)
    dataset = dataset.species_to_indices(species_order)
    dataset = dataset.shuffle()
    return dataset

dataset = load_ani_dataset("./ani_gdb_s01_to_s04.h5")
# Use dataset.split method to do split
train_data, val_data, test_data = dataset.split(0.8,0.1,0.1)

```

## Batching

```
In [5]: batch_size = 8192
# use dataset.collate(...).cache() method to do batching
train_data_loader = train_data.collate(batch_size).cache()
val_data_loader = val_data.collate(batch_size).cache()
test_data_loader = test_data.collate(batch_size).cache()

In [6]: total_train = sum(len(batch['energies']) for batch in train_data_loader)
total_val = sum(len(batch['energies']) for batch in val_data_loader)
total_test = sum(len(batch['energies']) for batch in test_data_loader)

print(f"Train: {total_train}")
print(f"Val: {total_val}")
print(f"Test: {total_test}")
print(f"Total: {total_train + total_val + total_test}")
```

```
Train: 691918
Val: 86489
Test: 86489
Total: 864896
```

## Torchani API

```
In [7]: class AtomicNet(nn.Module):
def __init__(self):
    super().__init__()
    self.layers = nn.Sequential(
        nn.Linear(384, 256),
        nn.BatchNorm1d(256),
        nn.ReLU(),
        nn.Dropout(0.15),
        nn.Linear(256, 192),
        nn.BatchNorm1d(192),
        nn.ReLU(),
        nn.Dropout(0.15),
        nn.Linear(192, 128),
        nn.BatchNorm1d(128),
        nn.ReLU(),
        nn.Dropout(0.15),
        nn.Linear(128, 1)
    )

def forward(self, x):
    return self.layers(x)

net_H = AtomicNet()
net_C = AtomicNet()
```

```

net_N = AtomicNet()
net_O = AtomicNet()

# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks`b
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)

```

```

In [8]: class Bigger_AtomicNet(nn.Module):
        def __init__(self):
            super().__init__()
            self.layers = nn.Sequential(
                nn.Linear(384, 512),
                nn.BatchNorm1d(512),
                nn.ReLU(),
                nn.Dropout(0.15),
                nn.Linear(512, 256),
                nn.BatchNorm1d(256),
                nn.ReLU(),
                nn.Dropout(0.15),
                nn.Linear(256, 128),
                nn.BatchNorm1d(128),
                nn.ReLU(),
                nn.Dropout(0.15),
                nn.Linear(128, 1)
            )

        def forward(self, x):
            return self.layers(x)

```

```

In [9]: class Bigger_lowerDO_AtomicNet(nn.Module):
        def __init__(self):
            super().__init__()
            self.layers = nn.Sequential(
                nn.Linear(384, 512),
                nn.BatchNorm1d(512),
                nn.ReLU(),
                nn.Dropout(0.05),
                nn.Linear(512, 256),
                nn.BatchNorm1d(256),
                nn.ReLU(),
                nn.Dropout(0.05),
                nn.Linear(256, 128),
                nn.BatchNorm1d(128),

```

```

        nn.ReLU(),
        nn.Dropout(0.05),
        nn.Linear(128, 1)
    )

    def forward(self, x):
        return self.layers(x)

```

```

In [10]: class Bigger_noDO_AtomicNet(nn.Module):
        def __init__(self):
            super().__init__()
            self.layers = nn.Sequential(
                nn.Linear(384, 512),
                nn.BatchNorm1d(512),
                nn.ReLU(),
                nn.Linear(512, 256),
                nn.BatchNorm1d(256),
                nn.ReLU(),
                nn.Linear(256, 128),
                nn.BatchNorm1d(128),
                nn.ReLU(),
                nn.Linear(128, 1)
            )

        def forward(self, x):
            return self.layers(x)

```

```

In [10]: train_data_batch = next(iter(train_data_loader))

loss_func = nn.MSELoss()
species = train_data_batch['species'].to(device)
coords = train_data_batch['coordinates'].to(device)
true_energies = train_data_batch['energies'].to(device).float()
_, pred_energies = model((species, coords))
loss = loss_func(true_energies, pred_energies)
print(loss)
train_data_batch

```

```

tensor(2.4881, device='cuda:0', grad_fn=<MseLossBackward0>)

```

```

Out[10]: defaultdict(list,
    {'species': tensor([[ 1,  1,  1, ...,  0, -1, -1],
        [ 1,  1,  1, ..., -1, -1, -1],
        [ 1,  2,  1, ..., -1, -1, -1],
        ...,
        [ 1,  1,  1, ...,  0, -1, -1],
        [ 1,  1,  2, ..., -1, -1, -1],
        [ 3,  1,  1, ..., -1, -1, -1]]),
    'coordinates': tensor([[-0.9120,  1.3640,  0.0050],
        [ 0.3138,  0.3679, -0.0376],
        [ 0.3482, -0.9044,  0.7329],
        ...,
        [-0.6108, -1.2486, -1.1242],
        [ 0.0000,  0.0000,  0.0000],
        [ 0.0000,  0.0000,  0.0000]],
        [[-0.0654, -0.8613, -0.0047],
        [ 0.7938,  0.4076,  0.0034],
        [-0.7395,  0.4685,  0.0065],
        ...,
        [ 0.0000,  0.0000,  0.0000],
        [ 0.0000,  0.0000,  0.0000],
        [ 0.0000,  0.0000,  0.0000]],
        [[ 1.4732,  0.0194,  0.0759],
        [ 0.1632, -0.0509, -0.5472],
        [-0.9548, -0.7146,  0.1509],
        ...,
        [ 0.0000,  0.0000,  0.0000],
        [ 0.0000,  0.0000,  0.0000],
        [ 0.0000,  0.0000,  0.0000]],
        ...,
        [[-0.8867,  1.2779, -0.0021],
        [ 0.3324,  0.3691,  0.0112],
        [ 0.3514, -0.8869,  0.7582],
        ...,
        [-0.5475, -1.2799, -1.3002],
        [ 0.0000,  0.0000,  0.0000],
        [ 0.0000,  0.0000,  0.0000]],
        [[ 1.3832, -0.4572, -0.0199],
        [ 0.5448,  0.7394,  0.0201],
        [-0.7114,  0.6814,  0.0091],
        ...,
        [ 0.0000,  0.0000,  0.0000],

```

```

        [ 0.0000,  0.0000,  0.0000],
        [ 0.0000,  0.0000,  0.0000]],

        [[-1.5551, -0.5071, -0.0944],
         [-0.6139,  0.6448,  0.0429],
         [ 0.7753,  0.1398, -0.1084],
         ...,
         [ 0.0000,  0.0000,  0.0000],
         [ 0.0000,  0.0000,  0.0000],
         [ 0.0000,  0.0000,  0.0000]]]),
        'energies': tensor([-0.0073,  0.0026,  0.0173, ..., -0.0196,  0.0056,  0.0371],
                             dtype=torch.float64)})

```

```

In [11]: class ANITrainer:
    def __init__(self, model, learning_rate, epoch, l2):
        self.model = model

        num_params = sum(item.numel() for item in model.parameters())
        print(f"{model.__class__.__name__} - Number of parameters: {num_params}")

        self.optimizer = torch.optim.Adam(model.parameters(), learning_rate, weight_decay=l2)
        self.epoch = epoch
        self.scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(self.optimizer, mode = 'min', patience = 10, factor

    def train(self, train_data, val_data, early_stop=True, draw_curve=True):
        self.model.train()

        # init data loader
        print("Initialize training data...")
        train_data_loader = train_data

        # definition of loss function: MSE is a good choice!
        loss_func = nn.MSELoss()

        # record epoch losses
        train_loss_list = []
        val_loss_list = []
        lowest_val_loss = np.inf

        for i in tqdm(range(self.epoch), leave=True):
            self.model.train()
            train_epoch_loss = 0.0
            total_train_samples = 0

            for train_data_batch in train_data_loader:

                # compute energies
                species = train_data_batch['species'].to(device)

```

```

        coords = train_data_batch['coordinates'].to(device)
        true_E = train_data_batch['energies'].to(device).float()
        _, pred_E = self.model((species, coords))

        # compute loss
        batch_loss = loss_func(pred_E, true_E)

        # do a step
        self.optimizer.zero_grad()
        batch_loss.backward()
        torch.nn.utils.clip_grad_norm_(self.model.parameters(), max_norm = 1.0)
        self.optimizer.step()

        batch_importance = len(train_data_batch['energies'])
        train_epoch_loss += batch_loss.item() * batch_importance
        total_train_samples += batch_importance

        # use the self.evaluate to get loss on the validation set

        train_epoch_loss /= total_train_samples
        self.model.eval()
        val_epoch_loss = self.evaluate(val_data)

        train_loss_list.append(train_epoch_loss)
        val_loss_list.append(val_epoch_loss)

        self.scheduler.step(val_epoch_loss)

        if early_stop:
            if val_epoch_loss < lowest_val_loss:
                lowest_val_loss = val_epoch_loss
                weights = self.model.state_dict()

    if draw_curve:
        fig, ax = plt.subplots(1, 1, figsize=(5, 4), constrained_layout=True)
        ax.set_yscale("log")
        # Plot train loss and validation loss
        ax.plot(range(len(train_loss_list)), train_loss_list, label='Train')
        ax.plot(range(len(val_loss_list)), val_loss_list, label='Validation')
        ax.legend()
        ax.set_xlabel("Epoch")
        ax.set_ylabel("Loss")

    if early_stop:
        self.model.load_state_dict(weights)

    return train_loss_list, val_loss_list

```



```

def evaluate(self, data, draw_plot=False):
    self.model.eval()
    # init data loader
    data_loader = data

    # init loss function
    loss_func = nn.MSELoss()
    total_loss = 0.0
    total_samples = 0

    if draw_plot:
        true_energies_all = []
        pred_energies_all = []

    with torch.no_grad():
        for batch_data in data_loader:

            # compute energies
            species = batch_data['species'].to(device)
            coords = batch_data['coordinates'].to(device)
            true_E = batch_data['energies'].to(device).float()
            _, pred_E = self.model((species, coords))

            # compute loss
            batch_loss = loss_func(pred_E, true_E)

            batch_importance = len(batch_data['energies'])
            total_loss += batch_loss.item() * batch_importance
            total_samples += batch_importance

            if draw_plot:
                true_energies_all.append(true_E.detach().cpu().numpy().flatten())
                pred_energies_all.append(pred_E.detach().cpu().numpy().flatten())

    if draw_plot:
        true_energies_all = np.concatenate(true_energies_all)
        pred_energies_all = np.concatenate(pred_energies_all)
        # Report the mean absolute error
        # The unit of energies in the dataset is hartree
        # please convert it to kcal/mol when reporting the mean absolute error
        # 1 hartree = 627.5094738898777 kcal/mol
        # MAE = mean(|true - pred|)
        hartree2kcalmol = 627.5094738898777
        mae = np.mean(np.abs(true_energies_all - pred_energies_all)) * hartree2kcalmol
        rmse = np.sqrt(np.mean((true_energies_all - pred_energies_all) ** 2)) * hartree2kcalmol
        fig, ax = plt.subplots(1, 1, figsize=(5, 4), constrained_layout=True)

```

```

ax.scatter(true_energies_all, pred_energies_all, label=f"MAE: {mae:.2f} kcal/mol\nRMSE: {rmse:.2f} kcal/mol")
ax.set_xlabel("Ground Truth")
ax.set_ylabel("Predicted")
xmin, xmax = ax.get_xlim()
ymin, ymax = ax.get_ylim()
vmin, vmax = min(xmin, ymin), max(xmax, ymax)
ax.set_xlim(vmin, vmax)
ax.set_ylim(vmin, vmax)
ax.plot([vmin, vmax], [vmin, vmax], color='red')
ax.legend()

return total_loss / total_samples

```

Checkpoint 2 template logs the loss after each epoch, so having the x-axis as '# batch' is inaccurate. Loss for training and evaluate are divided by the total number of samples to normalize per sample size. That allows for the training and loss curves to be on the same scale and comparable.

```

In [12]: class lower_patience_ANITrainer:
    def __init__(self, model, learning_rate, epoch, l2):
        self.model = model

        num_params = sum(item.numel() for item in model.parameters())
        print(f"{model.__class__.__name__} - Number of parameters: {num_params}")

        self.optimizer = torch.optim.Adam(model.parameters(), learning_rate, weight_decay=l2)
        self.epoch = epoch
        self.scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(self.optimizer, mode = 'min', patience = 5, factor

    def train(self, train_data, val_data, early_stop=True, draw_curve=True):
        self.model.train()

        # init data loader
        print("Initialize training data...")
        train_data_loader = train_data

        # definition of loss function: MSE is a good choice!
        loss_func = nn.MSELoss()

        # record epoch losses
        train_loss_list = []
        val_loss_list = []
        lowest_val_loss = np.inf

        for i in tqdm(range(self.epoch), leave=True):
            self.model.train()

```

```

train_epoch_loss = 0.0
total_train_samples = 0

for train_data_batch in train_data_loader:

    # compute energies
    species = train_data_batch['species'].to(device)
    coords = train_data_batch['coordinates'].to(device)
    true_E = train_data_batch['energies'].to(device).float()
    _, pred_E = self.model((species, coords))

    # compute loss
    batch_loss = loss_func(pred_E, true_E)

    # do a step
    self.optimizer.zero_grad()
    batch_loss.backward()
    torch.nn.utils.clip_grad_norm_(self.model.parameters(), max_norm = 1.0)
    self.optimizer.step()

    batch_importance = len(train_data_batch['energies'])
    train_epoch_loss += batch_loss.item() * batch_importance
    total_train_samples += batch_importance

# use the self.evaluate to get loss on the validation set

train_epoch_loss /= total_train_samples
self.model.eval()
val_epoch_loss = self.evaluate(val_data)

train_loss_list.append(train_epoch_loss)
val_loss_list.append(val_epoch_loss)

self.scheduler.step(val_epoch_loss)

if early_stop:
    if val_epoch_loss < lowest_val_loss:
        lowest_val_loss = val_epoch_loss
        weights = self.model.state_dict()

if draw_curve:
    fig, ax = plt.subplots(1, 1, figsize=(5, 4), constrained_layout=True)
    ax.set_yscale("log")
    # Plot train loss and validation loss
    ax.plot(range(len(train_loss_list)), train_loss_list, label='Train')
    ax.plot(range(len(val_loss_list)), val_loss_list, label='Validation')
    ax.legend()

```

```

        ax.set_xlabel("Epoch")
        ax.set_ylabel("Loss")

    if early_stop:
        self.model.load_state_dict(weights)

    return train_loss_list, val_loss_list

def evaluate(self, data, draw_plot=False):
    self.model.eval()
    # init data loader
    data_loader = data

    # init loss function
    loss_func = nn.MSELoss()
    total_loss = 0.0
    total_samples = 0

    if draw_plot:
        true_energies_all = []
        pred_energies_all = []

    with torch.no_grad():
        for batch_data in data_loader:

            # compute energies
            species = batch_data['species'].to(device)
            coords = batch_data['coordinates'].to(device)
            true_E = batch_data['energies'].to(device).float()
            _, pred_E = self.model((species, coords))

            # compute loss
            batch_loss = loss_func(pred_E, true_E)


            batch_importance = len(batch_data['energies'])
            total_loss += batch_loss.item() * batch_importance
            total_samples += batch_importance

            if draw_plot:
                true_energies_all.append(true_E.detach().cpu().numpy().flatten())
                pred_energies_all.append(pred_E.detach().cpu().numpy().flatten())

    if draw_plot:
        true_energies_all = np.concatenate(true_energies_all)
        pred_energies_all = np.concatenate(pred_energies_all)
        # Report the mean absolute error
        # The unit of energies in the dataset is hartree


```





```
In [26]: trainer.evaluate(test_data_loader_demo, draw_plot = True)
```

```
Out[26]: 0.00016267432094451252
```




```
In [27]: net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()

# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
trainer_moredata = ANITrainer(model, learning_rate = 1e-4, epoch = 30, l2 = 1e-5)
train_loss, val_loss = trainer_moredata.train(train_data_loader, val_data_loader)
```

Sequential – Number of parameters: 695556

Initialize training data...

100% | 30/30 [00:56<00:00, 1.88s/it]




```
In [28]: trainer_moredata.evaluate(test_data_loader, draw_plot = True)
```

```
Out[28]: 0.0002686463669388924
```








```
In [30]: trainer_moredata_1.evaluate(test_data_loader, draw_plot = True)
```

```
Out[30]: 0.5152190905448766
```




```
In [31]: net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()

# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
trainer_moredata_2 = ANITrainer(model, learning_rate = 1e-5, epoch = 100, l2 = 1e-5)
train_loss, val_loss = trainer_moredata_2.train(train_data_loader, val_data_loader)
```

Sequential – Number of parameters: 695556


Initialize training data...

100% | 100/100 [03:07<00:00, 1.87s/it]



```
In [32]: trainer_moredata_2.evaluate(test_data_loader, draw_plot = True)
```

```
Out[32]: 0.0003164563611975233
```




```
In [33]: net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()

# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
trainer_moredata_3 = ANITrainer(model, learning_rate = 1e-4, epoch = 100, l2 = 1e-5)
train_loss, val_loss = trainer_moredata_3.train(train_data_loader, val_data_loader)
```

Sequential – Number of parameters: 695556


Initialize training data...

100% | 100/100 [03:04<00:00, 1.85s/it]



```
In [34]: trainer_moredata_3.evaluate(test_data_loader, draw_plot = True)
```

```
Out[34]: 4.89487025321561e-05
```




```
In [35]: net_H = AtomicNet()
net_C = AtomicNet()
net_N = AtomicNet()
net_O = AtomicNet()

# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
trainer_moredata_4 = ANITrainer(model, learning_rate = 1e-4, epoch = 200, l2 = 1e-5)
train_loss, val_loss = trainer_moredata_4.train(train_data_loader, val_data_loader)
```

Sequential – Number of parameters: 695556

Initialize training data...


100% | 200/200 [06:06<00:00, 1.83s/it]



```
In [36]: trainer_moredata_4.evaluate(test_data_loader, draw_plot = True)
```

```
Out[36]: 2.4287488911197822e-05
```






```
In [41]: net_H = Bigger_AtomicNet()
net_C = Bigger_AtomicNet()
net_N = Bigger_AtomicNet()
net_O = Bigger_AtomicNet()

# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
trainer_moredata_5 = ANITrainer(model, learning_rate = 1e-4, epoch = 200, l2 = 1e-5)
train_loss, val_loss = trainer_moredata_5.train(train_data_loader, val_data_loader)
```

Sequential – Number of parameters: 1453060


Initialize training data...

100% | 200/200 [07:34<00:00, 2.27s/it]



```
In [42]: trainer_moredata_5.evaluate(test_data_loader, draw_plot = True)
```

```
Out[42]: 1.2384741794498457e-05
```




```
In [43]: net_H = Bigger_lowerDO_AtomicNet()
net_C = Bigger_lowerDO_AtomicNet()
net_N = Bigger_lowerDO_AtomicNet()
net_O = Bigger_lowerDO_AtomicNet()

# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
trainer_moredata_6 = ANITrainer(model, learning_rate = 1e-4, epoch = 200, l2 = 1e-5)
train_loss, val_loss = trainer_moredata_6.train(train_data_loader, val_data_loader)
```

Sequential – Number of parameters: 1453060

Initialize training data...


100% | 200/200 [07:36<00:00, 2.28s/it]



```
In [44]: trainer_moredata_6.evaluate(test_data_loader, draw_plot = True)
```


```
Out[44]: 6.609294922307067e-06
```





```
In [52]: trainer_moredata_no_D0.evaluate(test_data_loader, draw_plot = True)
```

```
Out[52]: 0.00011139973513727834
```




```
In [54]: net_H = Bigger_lowerDO_AtomicNet()
net_C = Bigger_lowerDO_AtomicNet()
net_N = Bigger_lowerDO_AtomicNet()
net_O = Bigger_lowerDO_AtomicNet()

# ANI model requires a network for each atom type
# use torch.ANIModel() to compile atomic networks
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(
    aev_computer,
    ani_net
).to(device)
trainer_moredata_lower_P = lower_patience_ANITrainer(model, learning_rate = 1e-4, epoch = 200, l2 = 1e-5)
train_loss, val_loss = trainer_moredata_lower_P.train(train_data_loader, val_data_loader)
```

Sequential – Number of parameters: 1453060

Initialize training data...


100% | 200/200 [07:37<00:00, 2.29s/it]



```
In [55]: trainer_moredata_lower_P.evaluate(test_data_loader, draw_plot = True)
```

```
Out[55]: 3.369527535516908e-05
```





Checkpoint 3 write up:

I started off with 3 hidden layers with 256 nodes in the first hidden layer, 192 nodes in the second, and 128 in the third. The model achieved a test MAE of 2.42 kcal/mol. It was trained on a learning rate of  $1e-4$ , 200 epochs, and an L2 rate of  $1e-5$ . I began training with 30 epochs, a learning rate of  $1e-4$ , and an L2 rate of  $1e-5$ . I noticed that the MAE was 7.89 kcal/mol. Since the loss curves were still decreasing after 30 epochs, I decided to increase the number of epochs. With the same learning rate, but with 100 epochs, the MAE was 3.51 kcal/mol. At 100 epochs, I noticed that the loss curves continued to decrease. I decided to increase the number of epochs and therefore achieved a lower MAE after 200 epochs. I also experimented with smaller initial learning rates of  $1e-5$  and  $1e-8$ . This turned out poorly because the models could not make any significant progress towards minimizing the loss. I could tell because the curves continued to decrease with a steep slope. Especially with a learning rate of  $1e-8$ , the loss curves were very steep and were not able to flatten out after 100 epochs. I also put a reduce LR on plateau learning rate scheduler. This way, if the scheduler sees that there has been no improvement in loss minimization after 10 epochs, it will multiply the learning rate by 0.5. I also set a floor learning rate of  $1e-8$  to prevent having the learning rate so small that there is no meaningful progress towards optimization. I began with a L2 hyperparameter of  $1e-5$ . L2 regression penalizes larger weights. Increasing the L2 hyperparameter would help if the model is overfitting. However, since the loss curves continue to decrease quite nicely, increasing the hyperparameter could hurt the result. I also decided to use the default batch size given from checkpoint 1. The loss curves do not look like they are fluctuating and the the loss continues to decrease. The data loader also loads the correct batch size of data into the model training, so no need for a hyperparameter when I initiate ANITrainer.

I then decided to increase the size of the hidden layers to 512 nodes in the first hidden layer, 256 nodes in the second, and 128 in the third. As a note, both architectures had 384 nodes as an input at 1 node as output. With a larger capacity, the MAE dropped to 1.81 kcal/mol. I then wanted to see if I could drop the MAE even lower by lowering drop out. I lowered drop out from 0.15 to 0.05. The result was 1.17 kcal/mol. For both cuves at 200 epochs, they began to flatten out. I even tried removing drop out all together to get it even lower. However, that resulted in an MAE of 4.92 kcal/mol. With the same parameters as the model that produced the 1.17 kcal/mol, I wanted to lower the patience of the learning rate scheduler from 10 to 5. This will allow for more frequent updates when the during later epochs when the loss curve is mostly flattened out. However, we see that the loss curve flattened out so much after 100 epochs. This tells us that the learning rate got so small that any progress towards minimizing loss was not meaningful. The MAE turned out to be 2.68 kcal/mol. Although I was trying to get MAE to be below 1, as that is the typical allowed error for what is considered good for a realistic chemical prediction, 1.17 kcal/mol was the closest I got so far.

```
In [13]: def traditional_torchani_kfold_cv(
    k,
    full_dataset,
    model_builder,
    epochs,
    batch_size=8192,
    learning_rate=1e-4,
    l2=1e-5,
    draw_curve=True,
    verbose=True,
):

    """
    K-fold CV for torchani models.

    Parameters
    -----
    k : int
        Number of folds.
    full_dataset : torchani TransformableIterable
    model_builder : callable
        Returns a fresh model each fold
    epochs, batch_size, learning_rate, l2 : training hyperparameters
    """
    if draw_curve:
        n_row, n_col = int(np.ceil(k / 3)), 3
        fig, axes = plt.subplots(n_row, n_col, figsize=(5*n_col, 4*n_row), constrained_layout=True)
        axes = axes.flatten()

    train_mse_list, test_mse_list = [], []
    test_mae_list, test_rmse_list = [], []
    trainers = []
```

```

shuffled = full_dataset.shuffle()
fractions = [1.0 / k] * k
folds = shuffled.split(*fractions)

for i in range(k):
    if verbose:
        print(f"\n Fold {i+1}/{k} ")

    test_fold = folds[i]
    val_idx = (i + 1) % k
    val_fold = folds[val_idx]
    train_fold_list = [folds[j] for j in range(k) if j != i and j != val_idx]

    train_batches = []
    for f in train_fold_list:
        for batch in f.collate(batch_size).cache():
            train_batches.append(batch)

    val_loader = val_fold.collate(batch_size).cache()
    test_loader = test_fold.collate(batch_size).cache()

    model = model_builder()
    trainer = ANITrainer(model, learning_rate=learning_rate, epoch=epochs, l2=l2)

    train_losses, val_losses = trainer.train(
        train_batches, val_loader,
        early_stop=True, draw_curve=False
    )

    test_mse = trainer.evaluate(test_loader, draw_plot=False)
    train_mse_list.append(trainer.evaluate(train_batches, draw_plot=False))
    test_mse_list.append(test_mse)

    hartree2kcalmol = 627.5094738898777
    model.eval()
    true_all, pred_all = [], []
    with torch.no_grad():
        for batch in test_loader:
            species = batch['species'].to(device)
            coords = batch['coordinates'].to(device)
            true_E = batch['energies'].to(device).float()
            _, pred_E = model((species, coords))
            true_all.append(true_E.cpu().numpy().flatten())
            pred_all.append(pred_E.cpu().numpy().flatten())

    true_all = np.concatenate(true_all)
    pred_all = np.concatenate(pred_all)
    mae = np.mean(np.abs(true_all - pred_all)) * hartree2kcalmol

```

```

rmse = np.sqrt(np.mean((true_all - pred_all) ** 2)) * hartree2kcalmol
test_mae_list.append(mae)
test_rmse_list.append(rmse)

trainers.append(trainer)

if draw_curve:
    axes[i].set_yscale("log")
    axes[i].plot(range(len(train_losses)), train_losses, label='Train')
    axes[i].plot(range(len(val_losses)), val_losses, label='Validation')
    axes[i].set_xlabel('Epoch')
    axes[i].set_ylabel('Loss')
    axes[i].set_title(f"Fold {i+1}")
    axes[i].legend()

if verbose:
    print(f"Fold {i+1} - MAE: {mae:.3f} kcal/mol, RMSE: {rmse:.3f} kcal/mol")

if draw_curve:
    for j in range(k, len(axes)):
        axes[j].axis('off')

if verbose:
    print(f"\n Final K-fold results")
    print(f"Test MAE: {np.mean(test_mae_list):.3f} +/- {np.std(test_mae_list):.3f} kcal/mol")
    print(f"Test RMSE: {np.mean(test_rmse_list):.3f} +/- {np.std(test_rmse_list):.3f} kcal/mol")
    print(f"Test MSE: {np.mean(test_mse_list):.3e} +/- {np.std(test_mse_list):.3e}")
    print(f"\nPer-fold MAE: {[f'{m:.3f}' for m in test_mae_list]}")
    print(f"Per-fold RMSE: {[f'{r:.3f}' for r in test_rmse_list]}")

return trainers, test_mae_list, test_rmse_list

```

```

In [14]: def make_best_model():
    nets = [Bigger_lowerDO_AtomicNet() for _ in range(4)]
    ani_net = torchani.ANIModel(nets)
    return nn.Sequential(aev_computer, ani_net).to(device)

trainers_2, train_errs_2, test_errs_2 = traditional_torchani_kfold_cv(
    k=10,
    full_dataset=dataset,
    model_builder=make_best_model,
    epochs=200,
    batch_size=8192,
    learning_rate=1e-4,
    l2=1e-5,
)

```

Fold 1/10  
Sequential – Number of parameters: 1453060  
Initialize training data...  
100%|██████████| 200/200 [10:20<00:00, 3.10s/it]  
Fold 1 – MAE: 1.477 kcal/mol, RMSE: 1.893 kcal/mol

Fold 2/10  
Sequential – Number of parameters: 1453060  
Initialize training data...  
100%|██████████| 200/200 [09:48<00:00, 2.94s/it]  
Fold 2 – MAE: 1.329 kcal/mol, RMSE: 1.656 kcal/mol

Fold 3/10  
Sequential – Number of parameters: 1453060  
Initialize training data...  
100%|██████████| 200/200 [09:42<00:00, 2.91s/it]  
Fold 3 – MAE: 2.132 kcal/mol, RMSE: 2.647 kcal/mol

Fold 4/10  
Sequential – Number of parameters: 1453060  
Initialize training data...  
100%|██████████| 200/200 [09:43<00:00, 2.92s/it]  
Fold 4 – MAE: 1.929 kcal/mol, RMSE: 2.443 kcal/mol

Fold 5/10  
Sequential – Number of parameters: 1453060  
Initialize training data...  
100%|██████████| 200/200 [09:45<00:00, 2.93s/it]  
Fold 5 – MAE: 1.016 kcal/mol, RMSE: 1.393 kcal/mol

Fold 6/10  
Sequential – Number of parameters: 1453060  
Initialize training data...  
100%|██████████| 200/200 [09:44<00:00, 2.92s/it]  
Fold 6 – MAE: 2.620 kcal/mol, RMSE: 3.087 kcal/mol

Fold 7/10  
Sequential – Number of parameters: 1453060  
Initialize training data...  
100%|██████████| 200/200 [09:55<00:00, 2.98s/it]  
Fold 7 – MAE: 1.181 kcal/mol, RMSE: 1.676 kcal/mol

Fold 8/10  
Sequential – Number of parameters: 1453060  
Initialize training data...

100%|██████████| 200/200 [09:50<00:00, 2.95s/it]

Fold 8 – MAE: 2.011 kcal/mol, RMSE: 2.522 kcal/mol

Fold 9/10

Sequential – Number of parameters: 1453060

Initialize training data...

100%|██████████| 200/200 [09:48<00:00, 2.94s/it]

Fold 9 – MAE: 1.078 kcal/mol, RMSE: 1.519 kcal/mol

Fold 10/10

Sequential – Number of parameters: 1453060

Initialize training data...

100%|██████████| 200/200 [09:46<00:00, 2.93s/it]

Fold 10 – MAE: 1.973 kcal/mol, RMSE: 2.388 kcal/mol

Final K-fold results

Test MAE: 1.675 +/- 0.506 kcal/mol


Test RMSE: 2.123 +/- 0.539 kcal/mol

Test MSE: 1.218e-05 +/- 6.007e-06


Per-fold MAE: ['1.477', '1.329', '2.132', '1.929', '1.016', '2.620', '1.181', '2.011', '1.078', '1.973']

Per-fold RMSE: ['1.893', '1.656', '2.647', '2.443', '1.393', '3.087', '1.676', '2.522', '1.519', '2.388']


Fold 1




Fold 2




Fold 3




Fold 4




Fold 5




Fold 6




Fold 7



Fold 8




Fold 9




Fold 10








```
In [15]: for trainer_kfold in trainers_2:  
         trainer_kfold.evaluate(test_data_loader, draw_plot=True)
```































```
In [16]: #torch.save(trainers_2[4].model.state_dict(), 'best_model_fold5_MAE101.pt')
```

```
In [17]: net_H = Bigger_lowerDO_AtomicNet()
net_C = Bigger_lowerDO_AtomicNet()
net_N = Bigger_lowerDO_AtomicNet()
net_O = Bigger_lowerDO_AtomicNet()
ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])
model = nn.Sequential(aev_computer, ani_net).to(device)
model.load_state_dict(torch.load('best_model_fold5_MAE101.pt', map_location=device, weights_only=True))
trainer_reload = ANITrainer(model, learning_rate=1e-4, epoch=200, l2=1e-5)
```

Sequential – Number of parameters: 1453060

```
In [18]: model.eval()
hartree2kcalmol = 627.5094738898777
all_true, all_pred, all_has_N, all_has_O = [], [], [], []
with torch.no_grad():
    for batch in test_data_loader:
        species = batch['species'].to(device)
        coords = batch['coordinates'].to(device)
        true_E = batch['energies'].to(device).float()
        _, pred_E = model((species, coords))
```

```

species_np = species.cpu().numpy()
true_np    = true_E.cpu().numpy().flatten()
pred_np    = pred_E.cpu().numpy().flatten()
for i, mol in enumerate(species_np):
    mol = mol[mol >= 0]
    all_true.append(true_np[i])
    all_pred.append(pred_np[i])
    all_has_N.append(2 in mol)
    all_has_O.append(3 in mol)
all_true = np.array(all_true)
all_pred = np.array(all_pred)
all_has_N = np.array(all_has_N)
all_has_O = np.array(all_has_O)
errors    = (all_true - all_pred) * hartree2kcalmol
categories = {
    'Hydrogen and Carbon (H+C)': ~all_has_N & ~all_has_O,
    'H+C+Nitrogen':      all_has_N & ~all_has_O,
    'H+C+Oxygen':      ~all_has_N & all_has_O,
    'H+C+Nitrogen+Oxygen': all_has_N & all_has_O,
}
print(f"'Category':<30} {'Count':>6} {'MAE':>10} {'RMSE':>10}")
print("-" * 58)
for name, mask in categories.items():
    if mask.any():
        e = errors[mask]
        print(f"{name:<30} {mask.sum():>6} {np.mean(np.abs(e)):>10.2f} {np.sqrt(np.mean(e**2)):>10.2f}")





```

| Category                  | Count | MAE  | RMSE |
|---------------------------|-------|------|------|
| Hydrogen and Carbon (H+C) | 17639 | 0.84 | 1.24 |
| H+C+Nitrogen              | 32299 | 1.13 | 1.52 |
| H+C+Oxygen                | 25244 | 0.84 | 1.14 |
| H+C+Nitrogen+Oxygen       | 11307 | 1.28 | 1.64 |

```

In [19]: fig, axes = plt.subplots(2, 2, figsize=(10, 8), constrained_layout=True)
for ax, (name, mask) in zip(axes.flatten(), categories.items()):
    e = errors[mask]
    mae = np.mean(np.abs(e))
    rmse = np.sqrt(np.mean(e ** 2))
    std = np.std(e)
    ax.hist(e, bins=50, edgecolor='black', alpha=0.7,
            label=f"n={mask.sum()}\nMAE={mae:.2f} kcal/mol\nRMSE={rmse:.2f} kcal/mol\nSTD={std:.2f} kcal/mol")
    ax.axvline(0, color='red', linestyle='--')
    ax.set_title(name)
    ax.set_xlabel("Error (kcal/mol)")
    ax.set_ylabel("Count")
    ax.legend(fontsize=8)

```







```
In [20]: colors = {'Hydrogen and Carbon (H+C)': 'blue', 'H+C+Nitrogen': 'green', 'H+C+Oxygen': 'orange', 'H+C+Nitrogen+Oxygen':  
fig, axes = plt.subplots(2, 2, figsize=(10, 8), constrained_layout=True)
```

```

for ax, (name, mask) in zip(axes.flatten(), categories.items()):
    mae = np.mean(np.abs(all_true[mask] - all_pred[mask])) * hartree2kcalmol
    rmse = np.sqrt(np.mean((all_true[mask] - all_pred[mask])** 2)) * hartree2kcalmol
    ax.scatter(all_true[mask], all_pred[mask], s=2, alpha=0.5, color=colors[name],
               label=f"{name}\nMAE={mae:.2f} kcal/mol\nRMSE={rmse:.2f} kcal/mol")
    ax.set_xlabel("Ground Truth")
    ax.set_ylabel("Predicted")
    xmin, xmax = ax.get_xlim()
    ymin, ymax = ax.get_ylim()
    vmin, vmax = min(xmin, ymin), max(xmax, ymax)
    ax.set_xlim(vmin, vmax)
    ax.set_ylim(vmin, vmax)
    ax.plot([vmin, vmax], [vmin, vmax], color='red')
    ax.legend(markerscale=5)

```



I chose to do 10 fold cross validation. For 10-fold cross validation, the data was split into 80% training, 10% validation data, and 10% test data. That is, one fold must be reserved for early stopping and the learning rate scheduler, another as the held out testing set, and the other 8 for training. Both early stop and ReduceLROnPlateau need to monitor loss on the validation data to decide when to save the best weights and when to reduce the learning rate. However, I also wanted to evaluate on the original 10% test\_data from test\_data\_loader using the 80/10/10 split. The best model produced between the 10 folds of the 10-fold cross-validation achieved an MAE = 1.01 kcal/mol and an RMSE = 1.38 kcal/mol on the original held-out test set. The ANI paper had an RMSE of 1.3 kcal/mol.

To explain the gap in error, we must look at the training data. In the ANI paper, the scientist's dataset contained on roughly 17.2 million conformations from roughly 57,951 molecules. Their dataset contained subsets of molecules up to 8 heavy atoms. I trained on only up to 4 heavy atoms. My dataset contained 846,896 conformations. My dataset was roughly 5% the size of their dataset. With fewer molecules, the model is exposed to fewer types of bonds, functional groups, and structures during training. Thus having less information to predict energies of new molecules.

The model architecture in the ANI paper is 768 -> 128 -> 128 -> 64 -> 1. They had 124,033 parameters in their network. They also had a 768-element AEV using 32 radial symmetry functions that probe out from each atom at evenly spaced distances with a cut off of 4.6 angstroms. The angular symmetry function uses 8 distance shifts and 8 angle shifts with a cut off of 3.1 angstroms. Having the atoms H, C, N, and O, the radial part produces 128 elements ( $4 * 32$ ). The angular parts produces  $10 * 8 * 8 = 640$  elements, where  $(4 * 5)/2$  unique atom type pairs. This totals to a 768 element AEV. My AEV describes each atom's local environment with roughly half the resolution. My AEV uses 16 radial shifts with a 5.2 angstrom cutoff and 8 angle shifts up to 3.5 angstroms as its cutoff. There are 64 ( $4 * 16$ ) radial elements and 320 angular components ( $10 * 4 * 8$ ). My architecture is 384 -> 512 -> 256 -> 128 -> 1. My network contains roughly 363,265 parameters. The ANI paper uses Gaussian activation with an exponential cost function. They reported that this method gives 2-4x lower error than standard approaches. I used ReLU and an MSE loss function. The ANI paper reduced the learning rate 6 times, starting from 0.001. Each time reducing by an order of magnitude after the model has converged under its current learning rate. I took a different approach and used ReduceLROnPlateau with a factor of 0.5 and a patience of 10. That is, when the model sees that the validation loss had no meaningful decrease within 10 epochs, the learning rate is halved.

Overall, my model achieved element category MAEs that are near or below the 1 kcal/mol chemical accuracy threshold. The MAEs are H+C = 0.84 kcal/mol, H+C+N = 1.13 kcal/mol, H+C+O = 0.84 kcal/mol, and H+C+N+O = 1.28 kcal/mol.